

Churn Predictor App

Model Training:

Data Pre – Processing:

Clearing the data

The first step included looking for the presence of null values in the data. In this case we found no null values.

Then we looked for any duplicate rows. This was also absent.

Then after analysing the headings of the table we can conclude that name and customer ID were irrelevant to the detection of Churn so those two columns were dropped from the table.

Data Visualisation

For calculation purposes values of Gender and location we mapped to particular numbers using dictionary.

Using **Heat Map** we determined the correlation between various features of the dataset. From this we concluded that there is no two features with high correlation so there was no need to go for feature engineering.

To further analyse the skewness and distribution of the data we plotted gaussian plot over histogram of each plots.

For detecting the presence of outliers we plotted **Box and Whisker Plot**. In this particular case there were no outliers.

Checking for outliers

Even though in this case there were no outliers but still using mathematical formulas of inter-Quartile range we calculated for outliers. From this also we found that there were no outliers.

Feature Scaling

Some Features like total_usage may have large value which have more impact on the result than other features. In order to reduce this influence we used feature scaling.

Here we used **min max scaling**. So for each feature the values were between the range 0 to 1.

Model Building

Here we used four models:

- a) Logistic Regression
- b) Kneighbour
- c) Decision Tree
- d) Random Forest
- e) Linear Support Vector Machine
- f) Gradient Boosting
- g) Light BGM

We trained all the models then analyse their performances.

Using the cross validation method we found that the performance of various models were as follows:

- a) Logistic Regression – 0.5006
- b) Kneighbour – 0.5021
- c) Decision Tree – 0.5010
- d) Random Forest – 0.4973
- e) Linear Support Vector Machine – 0.5022
- f) Gradient Boosting – 0.5028
- g) Light BGM – 0.5042

So LightBGM performed better when considering cross validation.

Then for further performance analysis we used confusion matrix and AUC ROC curve.

Confusion matrix of varios models:

a) Logistic Regression

```
LR
```

```
True Positive : 2556
True Negative : 7526
False Positive: 7424
False Negative: 2494
```

	precision	recall	f1-score	support
0.0	0.75	0.50	0.60	14950
1.0	0.26	0.51	0.34	5050
accuracy			0.50	20000
macro avg	0.50	0.50	0.47	20000
weighted avg	0.63	0.50	0.54	20000

b) Kneighbour

```
True Positive : 4909
True Negative : 5190
False Positive: 5071
False Negative: 4830
```

	precision	recall	f1-score	support
0.0	0.52	0.51	0.51	10261
1.0	0.49	0.50	0.50	9739
accuracy			0.50	20000
macro avg	0.50	0.50	0.50	20000
weighted avg	0.51	0.50	0.51	20000

c) Decision Tree

```
CART
```

```
True Positive : 4973
True Negative : 5086
False Positive: 5007
False Negative: 4934
```

	precision	recall	f1-score	support
0.0	0.51	0.50	0.51	10093
1.0	0.50	0.50	0.50	9907
accuracy			0.50	20000
macro avg	0.50	0.50	0.50	20000
weighted avg	0.50	0.50	0.50	20000

d) Random Forest

```
RF

True Positive : 4672
True Negative : 5318
False Positive: 5308
False Negative: 4702

Classification Report
              precision    recall  f1-score   support

     0.0         0.53      0.50      0.52     10626
     1.0         0.47      0.50      0.48      9374

 accuracy          0.50
 macro avg         0.50      0.50      0.50
weighted avg         0.50      0.50      0.50
```

e) Linear Support Vector Machine

```
SVR

True Positive : 9979
True Negative : 0
False Positive: 1
False Negative: 10020

Classification Report
              precision    recall  f1-score   support

     0.0         0.00      0.00      0.00         1
     1.0         1.00      0.50      0.67     19999

 accuracy          0.50
 macro avg         0.50      0.25      0.33
weighted avg         1.00      0.50      0.67
```

f) Gradient Boosting

```
GB

True Positive : 4045
True Negative : 5916
False Positive: 5935
False Negative: 4104

Classification Report
              precision    recall  f1-score   support

     0.0         0.59      0.50      0.54     11851
     1.0         0.41      0.50      0.45      8149

 accuracy          0.50
 macro avg         0.50      0.50      0.49
weighted avg         0.51      0.50      0.50
```

g) Light BGM

```
LightGBM
```

True Positive :	4251			
True Negative :	5765			
False Positive:	5729			
False Negative:	4255			

Classification Report				
	precision	recall	f1-score	support
0.0	0.58	0.50	0.54	11494
1.0	0.43	0.50	0.46	8506
accuracy			0.50	20000
macro avg	0.50	0.50	0.50	20000
weighted avg	0.51	0.50	0.50	20000

AUC ROC for all the models were:

- a) Logistic Regression – 0.5034
- b) Kneighbour - 0.5016
- c) Decision Tree – 0.5029
- d) Random Forest – 0.4991
- e) Linear Support Vector Machine – 0.5000
- f) Gradient Boosting – 0.4998
- g) Light BGM – 0.5043

So after analysing all the models based on performance we decided to deploy LightBGM model for the work.

Model Tunning

By using a max_depth of 6 and learning rate of 0.01 we tunned the model and save it in the system for deployment.

Model Deployment:

For model deploement we used Flusk , HTML , CSS.

App.py contains the code for running the app on the local host. This renders HTML which contains a website for form using this form we get data from users and based on that our model predicts the output.

App Performance:

Churn is 0

Age
Gender
Location
Subscription Length (in months)
Monthly Bill
Total Usage (in GB)

Submit form

Age: 63,
Gender: Male,
Location: Los Angeles,
Subscription Length:17,
Monthly Bill: 73.36,
Total_Usage: 236

Churn is 1

Age
Gender
Location
Subscription Length (in months)
Monthly Bill
Total Usage (in GB)

Submit form

Age: 20,
Gender: Female,
Location: Miami,
Subscription Length:10,
Monthly Bill: 42.45,
Total_Usage: 150